

Task 6 - Machine Learning 3

Overview

Having completed Task 5 and advanced the project codebase to version v0.5, I was assigned the task of developing an ensemble approach for combining multiple machine learning and statistical analysis methods for improving the prediction quality of the existing version of the project, which is the implementation of a single model of Long Short-Term Memory (LSTM).

As introduced in the task brief, the basic idea supporting ensemble methods is to combine multiple learning algorithms to obtain better predictive performance that could be obtained from any of the single learning algorithms alone. In this task, I would have to develop an ensemble modelling approach consisting of at least two models, autoregressive integrated moving average (ARIMA) or seasonal ARIMA (SARIMA), and the existing DL model using LSTM, as well as experiment with different ensemble models (e.g., ARIMA/SARIMA/Random Forest/LSTM/RNN/GRU) and with different hyperparameters configurations.

SARIMA

SARIMA is a forecasting method for univariate time series data that works well with data having a seasonal component, supporting both autoregressive and moving average elements, while the integrated element of differencing allowing the method to support non-stationary time series data.

The seasonality of a time series refers to recurring and regular patterns at specific regular intervals less than a year (i.e., weekly, monthly, or quarterly). It may be caused by out-of-scope factors belonging to the natural and societal world (e.g., weather, vacation, holidays). As stock data is recorded on a daily basis (excepting non-working days of Saturdays and Sundays), and is strongly influenced by the events in the economics and politics world, it has a great likelihood to express seasonality.




Figure 1: Example of seasonal time series data.

To kick off implementations of SARIMA, I added the `pmdarima` package to the `requirements.txt` file of the project, and create a new virtual environment. Here is the point where I encountered an error of inability to create a new virtual environment, because of the incompatibility of the `pmdarima` package with Python 3.13. As a result, I downgrade Python by installing Python 3.12, and created a new virtual environment using the earlier version. The process turned out to be successful.

```

requirements.txt
1 scikit-learn
2 yfinance
3 pandas
4 numpy
5 mplfinance
6 seaborn
7 tensorflow
8 pmdarima

```

Figure 2: All required packages of the project, up until this point of implementing SARIMA.

As mentioned in the attached blog post, there are a total of 7 parameters requiring configuration for the SARIMA model:

- p : Trend autoregression order.
- d : Trend difference order.
- q : Trend moving average order.
- P : Seasonal autoregressive order.
- D : Seasonal difference order.
- Q : Seasonal moving average order.
- m : The number of time steps for a single seasonal period.

The optimal values of each parameter could be found by doing a grid search fitting a pre-specified range of parameter values. The set of parameter values with the lowest AIC (aka. Akaike information criterion), which is an estimator of prediction error, would be chosen. The following code snippet load the data needed and perform the aforementioned process, after having imported the necessary modules and functions, and followed the blog post's procedure.

```

sarima.py > ...
1 import pmdarima as pm
2 from stock_prediction import load_data
3 from parameters import *
4
5 # Load the data.
6 data = load_data(
7     ticker=TICKER, start_date=START_DATE, end_date=END_DATE,
8     test_size=TEST_SIZE,
9     store_locally=STORE_LOCALLY, load_locally=LOAD_LOCALLY,
10    local_data_path=LOCAL_DATA_PATH,
11    scale_data=SCALE_DATA,
12    split_by_date=SPLIT_BY_DATE,
13    history_days=HISTORY_DAYS, predict_days=PREDICT_DAYS,
14    feature_columns=FEATURE_COLUMNS,
15 )
16
17 # Fitting a stepwise model to find the best parameters for SARIMA.
18 stepwise_fit = pm.auto_arima(
19     data['df']['Close'], start_p=1, start_q=1, max_p=3, max_d=2, max_q=3, m=7,
20     start_P=0, start_Q=0, max_P=3, max_D=3, max_Q=3, seasonal=True, trace=True,
21     error_action='ignore',
22     suppress_warnings=True,
23     stepwise=True,
24 )
25 print(stepwise_fit.summary())

```

Figure 3: The code to load the data, and perform grid search for the optimal set of parameters for the SARIMA model.

When trying to run this, I encountered a `ValueError` indicating a possible binary incompatibility of the `pmdarima` module towards the `numpy` one.

```
(.venv) PS E:\kel-25-vin\Swinburne\UG\2025 - Semester 2\COS30018 - Intelligent Systems\Assignment Project\FinTech101> & "E:/kel-25-vin/Swinburne/UG/2025 - Semester 2/COS30018 - Intelligent Systems/Assignment Project/FinTech101/.venv/Scripts/python.exe" "e:/kel-25-vin/Swinburne/UG/2025 - Semester 2/COS30018 - Intelligent Systems/Assignment Project/FinTech101/sarima.py"
--> from ..utils.array import c, check_endog
   file "E:/kel-25-vin/Swinburne/UG/2025 - Semester 2/COS30018 - Intelligent Systems/Assignment Project/FinTech101/.venv/lib/site-packages/pmdarima/utils/_init_.py", line 5, in <module>
     from .array import *
   File "E:/kel-25-vin/Swinburne/UG/2025 - Semester 2/COS30018 - Intelligent Systems/Assignment Project/FinTech101/.venv/lib/site-packages/pmdarima/utils/array.py", line 13, in <module>
     from .array import C_intgrt_vec
   File "E:/kel-25-vin/Swinburne/UG/2025 - Semester 2/COS30018 - Intelligent Systems/Assignment Project/FinTech101/.venv/lib/site-packages/pmdarima/utils/array.py", line 1, in init pmdarima.utils.array
ValueError: numpy.dtype size changed, may indicate binary incompatibility. Expected 96 from C header, got 88 from PyObject
```

Figure 4: The `ValueError` suggesting incompatibility of `pmdarima` with `numpy`.

After doing some research on the Internet, I found out that the issue here is somewhat similar to the issue making the virtual environment unable to be created. It's the fact that `pmdarima` is only compatible with `numpy` version 1.x, because the model was compiled using that version, not the version I was using, `numpy` 2.x. Consequently, I had to downgrade `numpy` to 1.x by running a simple command, and the code worked perfectly.

```
(.venv) PS E:\kel-25-vin\Swinburne\UG\2025 - Semester 2\COS30018 - Intelligent Systems\Assignment Project\FinTech101> python -m pip install "numpy<2"
Collecting numpy<2
  Downloading numpy-1.26.4-cp312-cp312-win_amd64.whl.metadata (61 kB)
  Downloading numpy-1.26.4-cp312-cp312-win_amd64.whl (15.5 MB)
    15.5/15.5 MB 10.7 MB/s 0:00:01
Installing collected packages: numpy
  Attempting uninstall: numpy
    Found existing installation: numpy 2.3.3
    Uninstalling numpy-2.3.3:
      Successfully uninstalled numpy-2.3.3
Successfully installed numpy-1.26.4
```

Figure 5: The command used to downgrade `numpy`.

```
Best model: ARIMA(0,1,0)(0,0,0)[7] intercept
Total fit time: 1.027 seconds
SARIMAX Results
=====
Dep. Variable: y No. Observations: 1412
Model: SARIMAX(0, 1, 0) Log Likelihood: -2429.916
Date: Sun, 12 Oct 2025 AIC: 4863.832
Time: 22:18:34 BIC: 4874.336
Sample: 0 HQIC: 4867.757
- 1412
Covariance Type: opg
=====
            coef  std err      z    P>|z|    [0.025    0.975]
-----
intercept  0.0789  0.037    2.117    0.034    0.006    0.152
sigma2    1.8338  0.037   49.315    0.000    1.761    1.907
=====
Ljung-Box (L1) (Q): 0.65  Jarque-Bera (JB): 1795.25
Prob(Q): 0.42  Prob(JB): 0.00
Heteroskedasticity (H): 1.96  Skew: -0.68
Prob(H) (two-sided): 0.00  Kurtosis: 8.36
=====
Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

Figure 6: The result of the stepwise grid search process.

With that done, I proceeded to make predictions with the result model. The predictions are made for the following 30 days of the input data.

```
# Fitting the SARIMA model with the best parameters found.
SARIMA_Forecast = pd.DataFrame(stepwise_fit.predict(n_periods=30))
SARIMA_Forecast.columns = ['SARIMA_Forecast']
SARIMA_OUTPUT = SARIMA_Forecast.reset_index()['SARIMA_Forecast']
print(SARIMA_OUTPUT)
```

Figure 7: The code to make prediction using the output SARIMAX model.

0	175.284570
1	175.363422
2	175.442273
3	175.521125
4	175.599976
5	175.678828
6	175.757679
7	175.836531
8	175.915382
9	175.994234
10	176.073085
11	176.151937
12	176.230788
13	176.309639
14	176.388491
15	176.467342
16	176.546194
17	176.625045
18	176.703897
19	176.782748
20	176.861600
21	176.940451
22	177.019303
23	177.098154
24	177.177006
25	177.255857
26	177.334709
27	177.413560
28	177.492411
29	177.571263

Name: SARIMA_Forecast, dtype: float64

Figure 8: The predictions made by the SARIMAX model.

Ensemble with LSTM

First, I made some modifications with the LSTM model's training process, by adding a `ModelCheckpoint` callback which allow me to retrieve and save the best model over the whole period of training. After training, the system would load the best model saved, and make predictions on the testing data. I also made a change which would help display the final predictions in the format of a DataFrame.

```

# Callbacks.
checkpointer = ModelCheckpoint(filepath=f'models/{MODEL_NAME}.keras', save_best_only=True, verbose=1)

# Train the model.
history = model.fit(data["X_train"], data["y_train"],
                      batch_size=BATCH_SIZE, epochs=EPOCHS,
                      validation_data=(data["X_test"], data["y_test"]),
                      callbacks=[checkpointer], verbose=1)

# Predict the future stock prices.
best_model = load_model(f'models/{MODEL_NAME}.keras', compile=False)
predicted_price = predict(best_model, data)
LSTM_OUTPUT = pd.DataFrame(predicted_price, columns=['LSTM_Prediction'])['LSTM_Prediction']
print(LSTM_OUTPUT)

```

Figure 9: The modified version of the `train.py` module, which now includes a checkpoint callback for model persistence.

Then, I changed the parameter values of `HISTORY_DAYS` and `PREDICT_DAYS` to 365 and 30 respectively. This change indicates that the model now would use data from a year-long window to predict the price of roughly a month after.

```

HISTORY_DAYS = 365
PREDICT_DAYS = 30

```

Figure 10: New parameter configurations allowing the model to use a year of price data to predict the price of approximately a month later.

With these done, I proceeded to train the LSTM model for 100 epochs, and make predictions with the best model.

```

Epoch 100/100
13/13 0s 817ms/step - loss: 7.3877e-04 - mae: 0.0302
Epoch 100: val_loss did not improve from 0.00280
13/13 11s 890ms/step - loss: 7.4061e-04 - mae: 0.0306 - val_loss: 0.0034 - val_mae: 0.0645

```

Figure 11: Training done for the LSTM model.

	LSTM_Prediction
0	186.764542
1	187.241913
2	187.992020
3	189.144730
4	190.444504
5	191.531540
6	192.169540
7	192.365540
8	192.274673
9	191.904694
10	191.203857
11	190.304245
12	189.393433
13	188.583649
14	188.005295
15	187.574310
16	187.200836
17	186.922073
18	186.587585
19	186.395752
20	186.490677
21	186.576218
22	186.271805
23	185.573425
24	184.783905
25	183.920792
26	183.144287
27	182.495209
28	182.110306
29	182.047974

Figure 12: Predictions made by the best saved model.

Having gathered the predictions of LSTM and SARIMA models, I proceeded to compute and evaluate the ensemble predictions. This code snippet below from the newly created `ensemble.py` module downloads the new data for 30 following days, calculate the ensemble prediction values, output the DataFrame constructed by the two Series of predicted and true prices respectively, and eventually evaluate the predictions using the mean absolute percentage error function.

```
ensemble.py > ...
1 import yfinance as yf
2 from train import LSTM_OUTPUT
3 from sarima import SARIMA_OUTPUT
4 import pandas as pd
5 from sklearn.metrics import mean_absolute_percentage_error
6
7 # Download new stock data for the ensemble prediction (30 days).
8 new_data = yf.download("CBA.AX", start='2025-07-31', end='2025-09-11', multi_level_index=False)
9 new_close_prices = new_data.reset_index()['Close']
10
11 # Calculate the ensemble prediction by averaging LSTM and SARIMA outputs.
12 ensemble_output = (LSTM_OUTPUT + SARIMA_OUTPUT) / 2
13 ensemble_df = pd.DataFrame({
14     'Ensemble_Prediction': ensemble_output,
15     'True_Close_Price': new_close_prices.values,
16 })
17 print(ensemble_df)
18 print(mean_absolute_percentage_error(ensemble_df['True_Close_Price'], ensemble_df['Ensemble_Prediction']))
```

Figure 13: Calculating and evaluating the ensemble predictions.

The recorded MAPE value was **6.68%**, showing better predictions than using only the existing LSTM model.

	Ensemble_Prediction	True_Close_Price
0	181.024564	175.205719
1	181.302667	172.399048
2	181.717147	172.271011
3	182.332927	174.644379
4	183.022248	176.357941
5	183.605184	175.422379
6	183.963617	173.925476
7	184.101035	175.885239
8	184.095028	176.082199
9	183.949464	166.549332
10	183.638471	164.668365
11	183.228091	165.613770
12	182.812110	167.603073
13	182.446644	168.449997
14	182.196893	172.399994
15	182.020826	173.800003
16	181.873507	172.839996
17	181.773552	170.300003
18	181.645741	169.080002
19	181.589250	169.830002
20	181.676138	173.350006
21	181.758342	170.300003
22	181.645561	168.949997
23	181.335797	170.460007
24	180.980455	164.550003
25	180.588324	168.000000
26	180.239498	168.139999
27	179.954384	168.240005
28	179.801358	166.080002
29	179.809626	168.539993
0.06684639045096678		

Figure 14: Ensemble results showing performance improvements.

References

- Seasonality on Kaggle. <https://www.kaggle.com/code/ryanholbrook/seasonality>.
- The solution to the SARIMA's `ValueError` on Reddit.
https://www.reddit.com/r/pythonhelp/comments/1dtofq7/import_pmdarima_as_pm_keeps_giving_valueerror/.
- The attached blog post. <https://medium.com/analytics-vidhya/combining-time-series-analysis-with-artificial-intelligence-the-future-of-forecasting-5196f57db913>.